

Prague University of Economics and Business
Faculty of Informatics and Statistics



**Volume, Specificity, and Boilerplate:
A Decomposition of 10-K Risk-Factor
Language and Post-Filing Stock
Return Volatility**

Bachelor Thesis

Study programme: Data Analytics

Author: Šimon Slanský

Supervisor: prof. RNDr. Ing. Michal Černý, Ph.D.

Prague, January 2027

Abstract

The Item 1A risk-factor section of the U.S. 10-K is the subject of two parallel literatures whose predictions point in opposite directions. One, represented by Campbell et al. (2014), links a larger *volume* of risk discussion to higher subsequent volatility. The other, represented by Hope et al. (2016), links a higher *specificity* of risk discussion to lower volatility. This thesis takes the question through three layers. It first separates the volume measurement (the log word count of Item 1A) from the specificity measurement (the Loughran–McDonald risk-word density) and enters them jointly in a panel regression of post-filing log-volatility. It then decomposes the composite LM risk-density into its two sub-lists—uncertainty wording and litigious wording—each normalised by Item 1A length, and enters them separately. All specifications use two-digit SIC industry and fiscal-year fixed effects and firm-clustered standard errors. On an unbalanced panel of 535 large U.S. firms over fiscal years 2010–2024 (5,691 firm-years at the 30-day horizon), longer Item 1A sections predict higher post-filing volatility ($\hat{\beta} = +0.130$, $t = +7.70$), while denser Item 1A sections, holding length constant, predict *lower* post-filing volatility ($\hat{\beta} = -3.88$, $t = -2.74$). The LM decomposition reveals that the entire negative effect is carried by the LM-Litigious sub-list ($\hat{\beta}^{\text{lit}} = -6.16$, $t = -3.35$), while the LM-Uncertainty sub-list is statistically indistinguishable from zero. The pattern survives at every horizon from 5 to 365 days and on both COVID-era and pre-COVID sub-samples; we read it in light of the literature on standardised cautionary boilerplate (Beatty et al., 2019; Cazier et al., 2021; Nelson & Pritchard, 2007).

Keywords

textual analysis, 10-K filings, SEC EDGAR, volatility, risk factors, Item 1A, Loughran–McDonald, LM-Litigious word list, boilerplate disclosure, panel data, corporate disclosures

Contents

Introduction	5
1 Literature Review	7
2 Data	8
2.1 Selecting the Sample	8
2.2 Choosing the Filing Frequency — why annual 10-Ks?	8
2.3 XBRL Tag Consistency Dilemma	10
2.4 How do we extract text from 10-K filings?	11
2.4.1 Source of textual data	11
2.4.2 Section identification and exhibit fallback	11
2.4.3 Post-Processing	12
2.4.4 Coverage and Limitations	12
3 Methodology	14
3.1 Empirical strategy	14
3.2 Variables	14
3.2.1 Dependent variable: post-filing log-volatility	14
3.2.2 Financial controls	15
3.2.3 Text variables: Item 1A length and risk-word density	15
3.2.4 Peer-relative omission gap (H3)	16
3.3 Model specifications	17
3.4 Estimation and inference	18
4 Results	20
4.1 Descriptive statistics	20
4.2 Volume, specificity, and boilerplate at the 30-day horizon	21
4.3 The same channels at the 365-day horizon and across intermediate windows	24
4.4 Multi-horizon table	26
4.5 The boilerplate decomposition	27
4.6 Strategic silence and within-firm post-filing volatility	27
5 Robustness	29
5.1 Alternative post-filing windows	29
5.2 Pre-COVID and COVID-era sub-samples	29
5.3 Within-firm identification: firm fixed effects	31
5.4 What did not work	33
Conclusion	35
References	36

A	Variable Definitions and Data Pipeline	39
A.1	Filing Frequency Diagnostics	39
A.2	XBRL Tag Provenance	41
B	Documentation of AI Tool Usage	42

List of Abbreviations

CIK Central Index Key

DV Dependent Variable

EDGAR Electronic Data Gathering,
Analysis, and Retrieval

EPS Earnings Per Share

FE Fixed Effects

FY Fiscal Year

FYE Fiscal Year End

GLM Generalized Linear Model

L&M Loughran and McDonald

NLP Natural Language Processing

OCF Operating Cash Flow

OLS Ordinary Least Squares

ROA Return on Assets

SEC Securities and Exchange Commission

SIC Standard Industrial Classification

XBRL eXtensible Business Reporting
Language

Introduction

The Item 1A risk-factor section of the U.S. 10-K, mandatory since 2005, is the longest and most carefully written piece of risk-disclosure language in U.S. corporate reporting. Two parallel literatures argue that this section conveys economically meaningful information about firm-level risk, but they reach opposite predictions about its sign. Campbell et al. (2014) show that *larger* risk-factor sections, both in raw word count and in the count of risk-related words, are associated with higher subsequent return volatility, market beta, and bid-ask spreads: when managers write more about risk, they are signalling more risk. Hope et al. (2016) show that, for a given length, risk-factor sections that contain a higher proportion of *specific* content—names, numbers, geographies—are associated with stronger analyst-forecast revisions and lower implied cost of capital: more concrete risk language pins down exposures and reduces uncertainty.

Much of the empirical literature measures both predictions with a single statistic: the dictionary-based risk-word density of Item 1A. This is the share of the section’s words that match the Loughran and McDonald (2011) uncertainty and litigious word lists. But the ratio has section length in its denominator. The volume channel sits in the numerator and the specificity channel sits in the denominator, so the two effects are mixed into one number that has been hard to sign cleanly in the data.

This thesis takes a small methodological step with a real payoff. It splits the dictionary risk-density into its two parts—the natural log of the Item 1A word count, and the risk-word density itself—and enters them jointly in a panel regression of post-filing log-volatility. With both regressors in the equation, length captures the volume channel net of any change in boilerplate share, and density captures the specificity channel net of any change in volume. Both predictions can therefore be tested in the same regression, and the data can choose between them.

On an unbalanced panel of 535 large U.S. non-financial firms over fiscal years 2010–2024, the joint specification at the 30-day horizon produces both predictions: longer Item 1A sections predict higher post-filing volatility (the volume effect, H1), while denser Item 1A sections, holding length constant, predict lower volatility (the specificity effect, H2). Both signs and significance hold at every post-filing horizon between 5 and 365 days and on both pre-COVID and COVID-era sub-samples. The joint specification also shows that the dictionary-only estimate of the specificity channel is roughly twice as large as the true effect; the omitted volume channel loads onto density through the negative correlation between length and density.

A second-stage analysis looks inside the LM risk lexicon itself. The composite RiskDensity used for H2 sums two LM word lists with very different production technologies. Splitting it into LM-Uncertainty and LM-Litigious sub-densities shows that the whole negative effect comes from the LM-Litigious sub-list. At the 30-day horizon, holding length fixed, LitDensity has a coefficient of -6.16 ($t = -3.38$) on post-filing log-volatility, while UncDensity is small

and not statistically different from zero. The pattern holds at every horizon and on both COVID sub-samples. The reading is the one suggested by Beatty et al. (2019), Cazier et al. (2021), and Nelson and Pritchard (2007): litigious-list language is standardised cautionary boilerplate produced under the safe-harbour regime of the 1995 Private Securities Litigation Reform Act, and is typical of mature, legally supervised firms. The decomposition is treated as a refinement of H2 rather than as a separate hypothesis: it shows which part of the LM lexicon drives the H2 specificity effect.

A third hypothesis (H3) introduces a peer-relative measure of strategic silence motivated by the classical disclosure-equilibrium literature (Dye, 1985; Verrecchia, 1983). For each firm-year we build an *omission gap*: the share of TF-IDF disclosure mass written by the firm’s two-digit-SIC industry peers in the same fiscal year that the focal firm itself does not mention (Section 3.2.4). Holding Item 1A length, RiskDensity, and the H1 and H2 controls fixed, a higher omission gap is positively associated with post-filing log-volatility under firm and year fixed effects ($\hat{\beta}^{\text{om}} = +1.345$, $t = +2.93$ at the 30-day horizon). The result survives an industry-shuffle placebo, gets stronger with stricter length controls and a large-peer-cell restriction, and remains significant when the omission gap is lagged by one fiscal year. To our knowledge, the H3 evidence is the first panel test of strategic silence in the post-filing realised-volatility literature, with two qualifications discussed in Chapter 5.

The remainder of the thesis is organised as follows. Chapter 1 reviews the volume, specificity, boilerplate, and strategic-silence literatures and motivates the three hypotheses. Chapter 2 describes the SEC EDGAR data, the textual extraction pipeline, and the construction of the financial controls and the post-filing volatility windows. Chapter 3 formalises the joint specification, the LM-list decomposition, and the peer-relative omission specification. Chapter 4 reports the headline tests of H1, H2, and H3. Chapter 5 reports the multi-horizon evidence, pre/post-COVID splits, a within-firm fixed-effects specification, the H3 robustness battery (placebo, length-residualised, large-cell, decile-rank, and lead-lag specifications), and a transparent account of seven alternative specificity-channel candidates that were tested and rejected. The Conclusion summarises the contribution and discusses limitations and suggested extensions.

1. Literature Review

2. Data

2.1 Selecting the Sample

We start with the sample of 1,000 largest U.S. public companies by market capitalisation, taken from the SEC EDGAR company-tickers file, and apply three filters. First, we keep only entities the SEC classifies as “operating,” which removes exchange-traded funds, closed-end funds, blank-cheque companies, and shell entities. Second, we exclude firms in SIC codes 6000–6999 (finance, insurance, and real estate), following common sample-selection practice in prior accounting research (Francis et al., 2005; Sloan, 1996). Third, we require at least five years of readable XBRL data from 10-K filings and 90% valid stock-return coverage, so that there is enough within-firm variation for fixed-effects estimation.

Sample size. We deliberately focus on large, liquid firms rather than the full EDGAR universe for several reasons. First, our regressions use firm fixed effects with standard errors clustered by firm. As Petersen (2009) shows, the reliability of cluster-robust standard errors depends on the number of clusters—here 539—not on total observations, and 539 clusters are well above the threshold for reliable inference. Also, the textual-analysis part of the thesis requires downloading, parsing, and scoring every 10-K in the sample. Expanding to several thousand firms would greatly increase computation time. Lastly, smaller firms tend to have shorter XBRL histories (mandatory electronic filing was phased in by firm size), more tagging errors, and thinner trading, all of which add noise to financial ratios and return-based variables.

The sample size is in line with prior work combining textual and financial analysis of annual filings. Campbell et al. (2014) study about 9,000 firm-years of risk-factor disclosures from large U.S. firms; Loughran and McDonald (2011) use a broader set of roughly 50,000 firm-years, but their goal—validating a sentiment dictionary—calls for maximal cross-sectional coverage, whereas the joint volume-specificity test estimated here exploits within-industry-year variation across ~ 535 firms.

2.2 Choosing the Filing Frequency — why annual 10-Ks?

Before moving to variable construction, it is worth explaining why this study uses annual 10-K filings instead of quarterly 10-Q reports. Quarterly data would roughly triple the sample, but two practical problems—shown in Tables 2.1–2.2—ruled out that approach.

Financial data gaps in Q4. 10-Q filings cover only Q1–Q3, so any quarterly panel must reconstruct Q4 as the difference between the annual figure and the sum of the three preceding

quarters. Table 2.1 shows that this residual approach is not usable for operating cash flow: the three quarterly facts coexist in only 7% of firm-years, because most firms report cumulative year-to-date cash flows rather than standalone quarterly amounts. Operating cash flow is the primary input into our OCF-to-assets ratio, so a quarterly panel would lose the cash-flow-based robustness check at the source. Even for net income and operating income, where coverage exceeds 80%, the implied Q4 disagrees with the explicit Q4 by more than one percent of the annual total in about 10–14% of comparable firm-years.

Table 2.1: Q4 Reliability: Availability and Consistency of Fourth-Quarter XBRL Facts

Metric	Firm-Years	Q1 – Q3 (%)	Explicit Q4 (%)	Comparable	Match (%)
Net Income	451	80.9	43.0	172	89.5
Operating Income	391	79.0	22.0	72	86.1
Operating Cash Flow	462	7.4	3.2	0	—

Note: “Q1 – Q3” indicates the share of fiscal years where all three quarterly facts (Q1, Q2, Q3) exist in XBRL. “Explicit Q4” indicates the share with a standalone fourth-quarter fact (period ≈ 90 days ending at fiscal year-end). “Comparable” is the count of firm-years for which both an explicit Q4 fact and the implied residual are available. “Match” reports the share of those comparable observations where the two values agree within 1% of the annual figure. Sample: 30 non-financial firms, all available fiscal years.

Missing risk-factor text in 10-Q. SEC Regulation S-K only requires firms to update risk factors in quarterly filings when “material changes” have occurred since the last 10-K. In practice, only about half the 10-Q filings in our pilot sample (15 firms, 90 filings) contained a legitimate risk-factor section. The rest had only a reference back to the annual report (Table 2.2). This held across sectors—technology, healthcare, consumer goods, energy—with few exceptions. Constructing the year-on-year text variables that the headline specifications require is therefore not feasible at a quarterly frequency.

Chosen frequency. For these reasons—and following prior textual-analysis work based on 10-K filings (Blankespoor et al., 2020; Li, 2008; Loughran & McDonald, 2011)—we use an annual panel. The XBRL coverage comparison (Table A.1) and a consolidated summary (Table A.2) are in Appendix A.

Table 2.2: Textual Data Availability: 10-K vs. 10-Q Filings

	10-K	10-Q
Filings tested	44	90
<i>MD&A (Management’s Discussion and Analysis)</i>		
Section found (%)	93	100
Median word count	20,494	5,806
<i>Risk Factors</i>		
Substantive section found (%)	100	52
Boilerplate / reference only (%)	0	48
Median word count (substantive)	29,902	9,144

Note: “Found” requires >200 words extracted by regex-based section parser. “Substantive” for 10-Q excludes filings containing only a reference to the annual 10-K (e.g., “no material changes from those disclosed in our Annual Report”). MD&A corresponds to Item 7 in 10-K and Part I Item 2 in 10-Q; Risk Factors corresponds to Item 1A in 10-K and Part II Item 1A in 10-Q. Sample: 15 firms, most recent filings.

2.3 XBRL Tag Consistency Dilemma

A practical problem with the XBRL Company Facts API is that a single accounting concept can appear under different US-GAAP taxonomy tags. For example, operating cash flow may be reported as either `NetCashProvidedByUsedInOperatingActivities` (total, including discontinued operations) or `...ContinuingOperations` (continuing operations only). Net income may be tagged as `NetIncomeLoss`, `ProfitLoss`, or `IncomeLossFromContinuingOperations`. Stockholders’ equity can include or exclude the non-controlling interest. These tags overlap a lot but are not identical: mixing them all within a single firm’s time series introduces noise, while picking one tag globally would lose firms that consistently report under the other.

We handle this problem with a two-part *tag-harmonisation* process: (i) *equivalence groups* for tags that are economically interchangeable, and (ii) *per-firm tag locking* for tags that are not.

Equivalence groups. When two tags differ only in taxonomy version—not in their economic meaning—we let both enter a firm’s time series. Specifically, the two operating-cash-flow tags and the pair `NetIncomeLoss/ProfitLoss` are treated as interchangeable. The third net-income variant, `IncomeLossFromContinuingOperations`, is *not* included because it deliberately excludes discontinued operations and can differ materially.

To verify this classification, we queried the SEC EDGAR Company Facts API for all 550 sample firms and identified 6,113 firm-year periods in which both tags of a pair were reported simultaneously; the overlap diagnostic supports treating each pair as interchangeable (median ROA bias 0.033 pp; details in Appendix A, Table A.4). The few cases with meaningful

differences are handled by a priority rule that prefers the parent-company tag whenever both are reported.

Per-firm tag locking. For metrics where the alternative tags *are* economically distinct (stockholders’ equity with vs. without NCI; `IncomeLossFromContinuingOperations` vs. `NetIncomeLoss`), we lock each firm to the tag it uses in the most fiscal years and set values from other tags to missing. Combined with the equivalence groups, this raises the complete-case rate for the five regression controls from 78.9% to 92.7% of firm-year observations. About 4% of firms end up with a continuing-operations OCF tag paired with a total net-income tag, a residual mismatch we accept and clip via winsorisation, in line with comparable XBRL studies (Chychyla & Kogan, 2015).

Per-firm tag choices and the overlap diagnostic are reported in Appendix A (Table A.4).

2.4 How do we extract text from 10-K filings?

This study heavily depends on the textual sections of 10-K reports Item 1A *Risk Factors* and Item 7 *Management’s Discussion and Analysis of Financial Condition and Results of Operations* (MD&A). This section describes how we find and extract those sections from the raw HTML filings on SEC EDGAR.

2.4.1 Source of textual data

All filings come directly from the SEC EDGAR via the Submissions API, which links each company’s Central Index Key (CIK) to the primary documents of its 10-K filings. From the panel of about 6,700 firm-year observations, we drop 10-K/A amendments. These typically consist of a cover page and a restated exhibit and do not contain the full narrative sections of the original 10-K making them unusable for further analysis. Financial data for amended fiscal years are kept, since the XBRL pipeline already picks up the amended figures when available (Section 2.3). Our extraction thus targets only original 10-K filings.

2.4.2 Section identification and exhibit fallback

Unlike structured XBRL data, the textual parts of 10-K filings are buried inside HTML with no standardised tagging, mixed with table-of-contents entries, running headers, and in-line cross-references. We use a four-stage HTML parser that (i) detects candidate headers using regex patterns for the target “Item X” strings combined with formatting cues (bold, ALL-CAPS, ≥ 12 pt font), rejecting hyperlinks and headers surrounded by other Item references within 300 characters (table-of-contents signature); (ii) tags confirmed headers in place before flattening to plain text; (iii) picks the candidate followed by the longest stretch of

text before the next expected Item header; and (iv) bounds the section at the next Item header (Item 1B/1C/2 for Item 1A; Item 7A/8 for Item 7). Sections shorter than 200 words trigger a fallback that follows incorporation-by-reference stubs into the relevant exhibit (defaulting to Exhibit 13) and reruns the same identification procedure with title-only patterns (“Risk Factors”, “Management’s Discussion and Analysis”) as backup. Implementation in `scripts/text_parser.py`.

2.4.3 Post-Processing

After HTML-to-text conversion the parser strips three layout artefacts: running “*<page number> Table of Contents*” headers, standalone page numbers on otherwise empty lines, and non-breaking-space characters. The cleaning rules target only patterns that cannot be part of meaningful text and are applied after section boundaries have already been set.

2.4.4 Coverage and Limitations

The extraction pipeline was validated on two samples: the 50 most recent filings (fiscal year 2024, both sections recovered in all 50) and a stratified random sample of 100 filings covering 2010–2024 (97/100 successful, with three failures driven by PDF-only exhibits or the absence of any formatting cue for the section header). On the full panel we expect a coverage rate of roughly 97%, with missing values treated as missing in the textual variable construction.

LLM-assisted quality audit. To check that the extracted sections are actually the right ones, we ran an independent validation with a large language model (Google Gemini 2.5 Flash, temperature = 0). A stratified random sample of 100 filings was drawn (20 per five-year period over 2010–2024) and both Item 1A and Item 7 were evaluated for every filing, giving 200 section-level judgments. The use of LLMs as independent annotators in this role is supported by Gilardi et al. (2023); the sample size delivers Wilson 95 % confidence intervals of roughly ± 3 percentage points (Agresti & Coull, 1998; Krippendorff, 2019). The model graded each section on five criteria (correct section identity, start- and end-boundary accuracy, content quality, completeness) on a three-level scale (Pass / Minor / Fail). The detailed results are reported in Appendix A (Table B.1): on every individual criterion the lenient (Pass + Minor) pass rate is above 95 %, and the overall rate across all five criteria at once is 93.5 % (187/200; 95 % CI: 89.2–96.2 %). Failures cluster in older filings (5/40 in 2010–2012 vs. 1/40 in 2022–2024), reflecting the gradual standardisation of SEC filing HTML over the sample.

Relation to prior approaches. Our extraction strategy is inspired by Dyer et al. (2017), who document substantial evolution in 10-K disclosure—length, boilerplate, and structure—over time and the challenges this creates for large-sample studies. Compared with plain-text

methods used in earlier work (Campbell et al., 2014; Li, 2008; Loughran & McDonald, 2011), our procedure avoids false positives from TOC entries and inline cross-references. Compared with XBRL-tagged textual blocks (e.g. `us-gaap:RiskFactorsTextBlock`), which became available with Inline XBRL adoption, our method covers a wider time span: structured textual tags are inconsistently filled before 2018, while the HTML-based approach works across the full 2010–2024 window.

3. Methodology

This chapter explains how the data described in Chapter 2 are used to test the three hypotheses stated in the introduction. Section 3.1 gives an overview of the empirical strategy. Section 3.2 defines the regression variables and Section 3.3 formalises the joint specification used for H1 and H2, the LM-list-decomposition specification used as the second-stage test of H2, and the peripheral-silence specification used for H3. Section 3.4 describes estimation and inference. The robustness battery is reported in Chapter 5.

3.1 Empirical strategy

The thesis tests two competing predictions about the same section of the 10-K—Item 1A, the mandatory risk-factor disclosure—in a single regression specification. The volume channel (Campbell et al., 2014) predicts that longer risk-factor sections indicate more risk and therefore higher post-filing volatility; the specificity channel (Hope et al., 2016) predicts that, holding length constant, a higher share of generic risk vocabulary indicates a more boilerplate, less informative section and is associated with *lower* post-filing volatility. Because the dictionary risk-density commonly used in prior work has section length in its denominator, separating the two channels requires entering log word count and risk-word density jointly. As a second stage, the composite RiskDensity is decomposed into the LM uncertainty- and litigious-list densities to identify which sub-list carries the H2 effect.

The dependent variable is the natural log of post-filing annualised volatility, which symmetrises the right-skewed raw distribution and makes coefficients interpretable as semi-elasticities.

3.2 Variables

3.2.1 Dependent variable: post-filing log-volatility

The dependent variable is the natural logarithm of the annualised standard deviation of daily log returns computed over a post-filing window of length h that begins two trading days after the 10-K filing date:

$$\ln \sigma_{i,t+1}^{[h]} = \ln \left(\sqrt{252} \cdot \text{sd}(r_{i,d})_{d \in W_{i,t}^{[h]}} \right), \quad (3.1)$$

where $r_{i,d} = \ln(P_{i,d}/P_{i,d-1})$ is the daily log return for firm i and $W_{i,t}^{[h]}$ is the post-filing window of length h calendar days for filing year t . The two-day lag avoids the same-day price reaction to the filing; the $\sqrt{252}$ factor rescales the daily standard deviation to an annualised figure. The headline horizon is $h = 30$ days; robustness checks at $h \in \{5, 10, 90, 180, 365\}$ trace the temporal decay of the signal (Chapter 5).

3.2.2 Financial controls

The controls are taken directly from the firm’s annual XBRL filings (Section 2.3). Five regressors enter every specification: *lagged log-volatility* $\ln \sigma_{i,t}^{[h]}$ (the same measure as the dependent variable but for the previous filing year’s window of identical length), *firm size* $\ln(\text{Total Assets})$, *leverage* (total liabilities over total assets), *return on assets* (net income over total assets), and *asset growth* (year-on-year change in total assets divided by lagged total assets). All ratios are winsorised at the 1st and 99th percentiles within the regression sample to bound the influence of extreme values without discarding economically meaningful observations.

3.2.3 Text variables: Item 1A length and risk-word density

Two text variables operationalise the volume and specificity channels.

Item 1A length (H1). The volume channel is captured by the natural logarithm of the cleaned word count of Item 1A:

$$\text{Words}_{i,t}^{1A} = \text{number of tokens in Item 1A of firm } i\text{'s year-}t \text{ 10-K}, \quad (3.2)$$

$\ln(\text{Words}_{i,t}^{1A})$ enters the regression. The cleaning pipeline (Section 2.4.3) strips XBRL tags, footnote markers, and table cell contents before tokenisation, so the count reflects genuine narrative content rather than markup. Across the sample the median Item 1A contains roughly 8,800 words, with a long right tail (Section 4.1).

Item 1A risk-word density (H2) and its LM-list decomposition. The specificity channel is captured by the share of words in Item 1A that match the Loughran and McDonald (2011) uncertainty or litigious word lists:

$$\text{RiskDensity}_{i,t} = \frac{n_{i,t}^{\text{unc}} + n_{i,t}^{\text{lit}}}{\text{Words}_{i,t}^{1A}}, \quad (3.3)$$

where n^{unc} and n^{lit} are the per-section counts of uncertainty- and litigious-list words. Higher values indicate that a larger share of the section consists of generic risk vocabulary; lower values indicate a section in which the narrative is dominated by firm-specific or non-dictionary content—names, numbers, geographies, products—that the Hope et al. (2016) reading associates with greater specificity.

For the second-stage test of H2 the same composite is decomposed into its two LM sub-lists, each normalised by Item 1A length:

$$\text{UncDensity}_{i,t} = \frac{n_{i,t}^{\text{unc}}}{\text{Words}_{i,t}^{1A}}, \quad (3.4)$$

$$\text{LitDensity}_{i,t} = \frac{n_{i,t}^{\text{lit}}}{\text{Words}_{i,t}^{1A}}, \quad (3.5)$$

with the algebraic identity $\text{RiskDensity} = \text{UncDensity} + \text{LitDensity}$. The two sub-densities operationalise the conceptually distinct uncertainty and litigious constructs: UncDensity captures hedging vocabulary about what management does not know about its own business, while LitDensity captures the cautionary legal-process boilerplate that legal advisers insert under the safe-harbour regime of Nelson and Pritchard (2007).

3.2.4 Peer-relative omission gap (H3)

To test H3 we construct a peer-relative measure of strategic silence. For each fiscal year t we form a TF-IDF representation of every Item 1A in the panel, restricting the vocabulary to terms appearing in at least five filings overall ($\text{min_df} = 5$) and at most 95% of filings ($\text{max_df} = 0.95$); rows are ℓ_2 -normalised. Let $\mathbf{x}_{i,t}$ denote the TF-IDF row vector for firm i in year t . For each two-digit SIC industry j in year t with at least six panel filings, we compute the leave-one-out peer aggregate

$$\mathbf{p}_{i,t} = \sum_{k \in S_{j(i),t}, k \neq i} \mathbf{x}_{k,t}, \quad (3.6)$$

where $S_{j,t}$ is the set of panel firms in SIC2 industry j that filed in year t . The peer aggregate $\mathbf{p}_{i,t}$ summarises the topical content of what the focal firm's industry peers chose to write that year, with the focal firm itself excluded. The *omission gap* is the share of peer disclosure mass on terms the focal firm does not mention:

$$\text{OMISSION}_{i,t}^{1A} = 1 - \frac{\sum_{w: x_{i,t,w} > 0} p_{i,t,w}}{\sum_w p_{i,t,w}}, \quad (3.7)$$

where the sum in the numerator is over the vocabulary terms that appear with non-zero weight in the focal firm's Item 1A. By construction $\text{OMISSION}_{i,t}^{1A} \in [0,1]$. A value of zero means the focal firm mentions every term any peer mentions; a value close to one means the firm's own vocabulary covers a small share of what its peers chose to discuss. Firm-years in SIC2-year cells with fewer than five peers are dropped because the peer aggregate is too noisy to support inference.

The measure has two properties that make it complementary to the H1–H2 text variables. First, the peer aggregate $\mathbf{p}_{i,t}$ is a function of *other* firms' choices and varies non-trivially within firm-year as the composition and content of industry-year peers change over time. The within-firm standard deviation of OMISSION^{1A} is 0.054 in the regression sample, 46% of the cross-sectional standard deviation of 0.117, so the firm-fixed-effects test of H3 has real variance to exploit (in contrast to the LM sub-densities used in the H2 decomposition, which are largely absorbed by firm fixed effects; Section 5.3). Second, the measure is mechanically negatively correlated with Item 1A length: a longer Item 1A trivially mentions more peer-disclosed terms. The Pearson correlation between OMISSION^{1A} and $\ln(\text{Words}^{1A})$ is -0.85 in the regression sample. The H3 specification therefore enters $\ln(\text{Words}^{1A})$ as a control alongside OMISSION^{1A} , so that the omission coefficient is interpreted holding length fixed. A length-residualised version of the variable is reported as a robustness check in Section 5.3.

The peer-relative construction requires the SIC2 cell to be both the relevant comparison set and a non-degenerate one. SIC2 is the standard peer benchmark in the disclosure-similarity literature (Brown & Tucker, 2011; Dyer et al., 2017); the cell-size median is 35 peers ($P25 = 14$). Section 5.3 reports a placebo specification in which SIC2 cell labels are randomly permuted within fiscal year before recomputing the omission gap. The H3 specification replaces *asset growth* with two liquidity proxies—the *current ratio* and *operating cash flow over assets*—in the financial control vector. The substitution follows the disclosure-cost logic of Dye (1985) and Verrecchia (1983): the incentive to withhold an industry-relevant risk topic should be increasing in the firm’s proprietary or financing costs of disclosure, for which short-term liquidity is the standard empirical proxy. Asset growth, which captures the volume channel for H1 and H2, plays no analogous role in the silence equilibrium.

3.3 Model specifications

Let $\mathbf{X}_{i,t}^{\text{fin}}$ collect the four financial controls (size, leverage, ROA, asset growth) and let $\sigma_{i,t}^{[h]}$ denote the firm’s lagged volatility computed on the prior filing year’s window of the same length h .

Joint specification (H1 and H2).

$$\ln \sigma_{i,t+1}^{[h]} = \alpha + \phi \ln \sigma_{i,t}^{[h]} + \beta' \mathbf{X}_{i,t}^{\text{fin}} + \beta^{\text{len}} \cdot \ln(\text{Words}_{i,t}^{1A}) + \beta^{\text{dens}} \cdot \text{RiskDensity}_{i,t} + \mu_j + \tau_t + \varepsilon_{i,t}, \quad (3.8)$$

where μ_j is a two-digit SIC industry fixed effect and τ_t is a fiscal-year fixed effect. Equation 3.8 is estimated at the headline horizon $h = 30$ days for the main table (Section 4.2) and at $h \in \{5, 10, 30, 90, 180, 365\}$ for the multi-horizon table 4.5. H1 is supported if $\beta^{\text{len}} > 0$ and significant at the 5% level; H2 is supported if $\beta^{\text{dens}} < 0$ and significant at the 5% level. Both predictions must hold jointly in equation 3.8 for the volume-specificity decomposition to be considered confirmed.

LM-list decomposition specification (second-stage test of H2). For the second stage of H2 the composite RiskDensity in equation 3.8 is replaced by its two LM sub-densities entered jointly:

$$\ln \sigma_{i,t+1}^{[h]} = \alpha + \phi \ln \sigma_{i,t}^{[h]} + \beta' \mathbf{X}_{i,t}^{\text{fin}} + \beta^{\text{len}} \cdot \ln(\text{Words}_{i,t}^{1A}) + \beta^{\text{unc}} \cdot \text{UncDensity}_{i,t} + \beta^{\text{lit}} \cdot \text{LitDensity}_{i,t} + \mu_j + \tau_t + \varepsilon_{i,t} \quad (3.9)$$

By the algebraic identity $\text{RiskDensity} = \text{UncDensity} + \text{LitDensity}$, the decomposition (equation 3.9) nests the joint specification (equation 3.8) as the special case $\beta^{\text{unc}} = \beta^{\text{lit}}$. Equation 3.9 is estimated at the headline $h = 30$ days for column (5) of the headline table (Section 4.2) and across all six post-filing horizons in Table 4.5. The H2 decomposition is supported if $\beta^{\text{lit}} < 0$ at the 5% level with the sign and significance preserved at every horizon, while β^{unc} is small in absolute value and statistically indistinguishable from zero. The

Pearson correlation between UncDensity and LitDensity is essentially zero ($r \approx -0.05$ in the regression sample), so the two coefficients are separately identified.

Strategic-silence specification (H3). For H3 the joint specification of equation 3.8 is augmented with the peer-relative omission gap of equation 3.7:

$$\ln \sigma_{i,t+1}^{[h]} = \alpha + \phi \ln \sigma_{i,t}^{[h]} + \beta' \mathbf{X}_{i,t}^{\text{fin}} + \beta^{\text{len}} \cdot \ln(\text{Words}_{i,t}^{1A}) + \beta^{\text{dens}} \cdot \text{RiskDensity}_{i,t} + \beta^{\text{om}} \cdot \text{OMISSION}_{i,t}^{1A} + \mu_j + \tau_t \quad (3.10)$$

Equation 3.10 is estimated under both fixed-effect structures: the cross-sectional version with μ_j a SIC2 industry fixed effect (the headline H1+H2 specification) and the strict version with μ_j replaced by a firm fixed effect μ_i , which absorbs all time-invariant firm-level heterogeneity. The within-firm test is the more demanding of the two because it isolates the year-to-year variation in omission driven by changes in peer disclosure choices. Both specifications are estimated at every horizon $h \in \{5, 10, 30, 90, 180\}$ days¹ because the H3 prediction is about post-filing risk incorporation over the same multi-week windows as H1 and H2.

The sample sizes for the H3 regressions are 4,630 to 4,636 firm-years, smaller than the H1+H2 sample of 5,691 at the 30-day horizon because of the ≥ 5 peers-per-cell restriction needed to construct the peer aggregate.

3.4 Estimation and inference

Equation 3.8 is estimated by ordinary least squares with two-digit SIC industry and fiscal-year fixed effects absorbed using the absorbing-regression routine of the `linearmodels` package. Coefficients are therefore identified from *within-industry*, *within-year* variation, which removes the risk that the results are driven by either time trends in volatility (e.g. the 2008 and 2020 crises) or by persistent industry-level differences in risk vocabulary or risk levels.

Standard errors. Standard errors are two-way clustered by firm and fiscal year (Cameron et al., 2011; Petersen, 2009). Two-way clustering is necessary because (i) consecutive firm-year volatility windows mechanically overlap, inducing strong within-firm serial correlation, and (ii) calendar-year shocks (most visibly the 2020 COVID shock) generate strong within-year cross-sectional correlation that one-way firm clustering would miss.

Outliers. The financial ratios are winsorised at the 1st and 99th percentiles of the regression sample. The text variables are not winsorised: $\ln(\text{Words}^{1A})$ is naturally bounded by the log

¹The $h = 365$ day horizon is dropped from the H3 battery because the omission gap requires the firm-year to be linked to a SIC2 cell of at least five peers in the same fiscal year, and the post-filing 365-day window for fiscal years 2023–2024 extends past the data cut-off; the $h \in \{5, 10, 30, 90, 180\}$ ladder retains the full overlap with the H1+H2 horizon panel of Table 4.5.

transformation and RiskDensity is a ratio bounded between zero and one in practice (sample maximum ≈ 0.05).

Sample composition. The estimation sample at the 30-day horizon consists of 5,691 firm-years on 535 unique firms; the 365-day horizon retains 6,172 firm-years on the same set of firms.

4. Results

This chapter reports the empirical results. Section 4.1 summarises the regression sample. Section 4.2 presents the headline test of H1 and H2 at the 30-day horizon and introduces the LM-list decomposition column (the second-stage test of H2). Section 4.3 reproduces the same specification at the 365-day horizon to confirm that the results are not an artefact of the short-window choice. Section 4.4 runs the joint H1+H2 decomposition at every horizon between 5 and 365 days, and Section 4.5 discusses the LM-list decomposition evidence formally. Section 4.6 introduces the strategic-silence specification (H3) and reports the headline within-firm test, with full robustness deferred to Chapter 5.

4.1 Descriptive statistics

Table 4.1 reports summary statistics for the regression sample at the 30-day horizon. The dependent variable has mean $\ln \sigma_{i,t+1}^{[30]} \approx -1.34$ (SD ≈ 0.59), corresponding to a typical annualised post-filing volatility around 0.26. Firm size centres on \$11.7 billion in total assets, confirming the focus on large, liquid firms. The two text variables show economically meaningful spread: $\ln(\text{Words}^{1A})$ has a standard deviation of 0.64 (median Item 1A $\approx 8,800$ words), and the LM uncertainty and litigious sub-shares are essentially uncorrelated ($r \approx -0.05$), confirming that the H2 decomposition delivers two genuinely separate regressors.

Table 4.1: Summary Statistics

	Mean	SD	P25	Median	P75
$\ln(\sigma_{i,t+1})$	-1.344	0.588	-1.740	-1.410	-1.058
$\ln(\sigma_{i,t})$	-1.353	0.590	-1.749	-1.424	-1.065
$\ln(\text{Total Assets})$	23.182	1.452	22.191	23.206	24.236
Leverage	0.616	0.218	0.476	0.611	0.739
ROA	0.063	0.087	0.027	0.062	0.106
Asset Growth	0.113	0.243	0.001	0.059	0.142
$\ln(\text{Words}^{1A})$	9.026	0.637	8.639	9.079	9.418
RiskDensity	0.051	0.006	0.047	0.051	0.055
UncDensity	0.034	0.005	0.032	0.034	0.037
LitDensity	0.017	0.005	0.013	0.016	0.019
Observations	5,691				
Firms	535				

Note: Summary statistics for the regression sample at the 30-day horizon. $\sigma_{i,t+1}$ is the annualised standard deviation of daily log returns over the 30-day window starting two trading days after the filing date; $\ln(\sigma_{i,t+1})$ is its natural logarithm. Words^{1A} is the cleaned word count of Item 1A; RiskDensity is the share of LM uncertainty- and litigious-list words in Item 1A; UncDensity and LitDensity are the two sub-components separately. Financial ratios are winsorised at the 1st and 99th percentiles.

Table 4.2 reports pairwise Pearson correlations. Lagged log-volatility is the strongest single predictor of log-volatility ($r = 0.35$). Larger firms have lower log-volatility ($r \approx -0.19$) and higher profitability also tracks lower log-volatility ($r \approx -0.23$), in line with the well-known size–risk and earnings-quality relationships. The two text variables are correlated at $r = -0.30$: longer Item 1A sections tend to have lower risk-word density, consistent with their narrative being padded by firm-specific text rather than by extra generic risk vocabulary. The negative correlation is moderate and leaves both signals separately identifiable in the joint specification.

Table 4.2: Pearson Correlation Matrix

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
(1) $\ln(\sigma_{i,t+1})$	1							
(2) $\ln(\sigma_{i,t})$	0.35***	1						
(3) $\ln(\text{Total Assets})$	-0.19***	-0.20***	1					
(4) Leverage	-0.06***	-0.07***	0.21***	1				
(5) ROA	-0.23***	-0.25***	0.05***	-0.06***	1			
(6) Asset Growth	0.16***	0.09***	-0.12***	-0.12***	-0.04***	1		
(7) $\ln(\text{Words}^{1A})$	0.31***	0.31***	-0.09***	-0.04***	-0.27***	0.17***	1	
(8) RiskDensity	-0.11***	-0.12***	-0.01	-0.01	0.06***	-0.08***	-0.30***	1
(9) UncDensity	-0.09***	-0.08***	-0.18***	-0.06***	0.14***	-0.05***	-0.43***	0.71***
(10) LitDensity	-0.06***	-0.08***	0.18***	0.04***	-0.06***	-0.06***	0.03**	0.66***

Note: Pairwise Pearson correlation coefficients on the 30-day-horizon regression sample. ***, **, * denote significance at the 1%, 5%, and 10% levels.

4.2 Volume, specificity, and boilerplate at the 30-day horizon

Table 4.3 reports the headline test of H1 and H2, together with the LM-list decomposition column. Column (1) is the financial baseline; Column (2) adds Item 1A length in logs (the volume channel); Column (3) adds the composite risk-word density (the specificity channel); Column (4) is the joint specification of equation 3.8 that nests H1 and H2; Column (5) is the LM-list decomposition specification of equation 3.9, which replaces RiskDensity by its two LM sub-components.

The baseline (Column 1). Lagged log-volatility, firm size (negative), ROA (negative), and asset growth (positive) enter with the expected signs and conventional significance; leverage is statistically indistinguishable from zero, as in much of the volatility literature (Campbell et al., 2014). The baseline already explains 49% of the in-sample variation.

The volume channel (Columns 2, 4, and 5). Adding $\ln(\text{Words}^{1A})$ alone (Column 2) raises adjusted R^2 to 0.505 and yields a length coefficient of +0.138 ($t = +8.43$). In the joint

Table 4.3: Item 1A Length, Risk-Word Density, and the Boilerplate Decomposition (30 days window)

	(1)	(2)	(3)	(4)	(5)
$\ln(\sigma_{i,t})$	+0.1695** (+2.12)	+0.1368* (+1.80)	+0.1552** (+1.96)	+0.1336* (+1.76)	+0.1321* (+1.75)
$\ln(\text{Total Assets})$	-0.0700*** (-5.81)	-0.0671*** (-5.91)	-0.0695*** (-5.99)	-0.0672*** (-5.98)	-0.0645*** (-5.69)
Leverage	-0.0198 (-0.56)	-0.0026 (-0.07)	-0.0249 (-0.72)	-0.0072 (-0.21)	-0.0082 (-0.24)
ROA	-1.0678*** (-9.72)	-0.8982*** (-8.68)	-1.0689*** (-9.80)	-0.9189*** (-8.80)	-0.9250*** (-8.77)
Asset Growth	+0.2460*** (+4.34)	+0.2033*** (+3.80)	+0.2313*** (+4.21)	+0.2011*** (+3.78)	+0.1982*** (+3.72)
$\ln(\text{Words}^{1A})$		+0.1375*** (+8.43)		+0.1212*** (+7.64)	+0.1298*** (+7.16)
RiskDensity			-7.7820*** (-5.40)	-3.8769*** (-2.68)	
UncDensity					-1.2779 (-0.62)
LitDensity					-6.1591*** (-3.38)
Industry FE	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes
Observations	5,691	5,691	5,691	5,691	5,691
Firms	535	535	535	535	535
Adj. R^2	0.492	0.505	0.497	0.506	0.507

Note: The dependent variable is $\ln(\sigma_{i,t+1})$, the natural log of annualised post-filing volatility over a 30 days window starting two trading days after the 10-K filing date. Column (1) reports the financial baseline; column (2) adds the log word count of Item 1A (H1); column (3) adds the composite LM risk-word density (H2); column (4) adds H1 and H2 jointly; column (5) replaces the H2 composite by its LM-list decomposition into the LM-Uncertainty and LM-Litigious densities of Item 1A. All specifications include two-digit SIC industry and fiscal-year fixed effects. t -statistics in parentheses use two-way (firm and fiscal-year) clustered standard errors. ***, **, * denote significance at the 1%, 5%, and 10% levels.

specification (Column 4), where the composite risk-word density is controlled for, the length coefficient is $+0.121$ ($t = +7.64$). In the LM-list decomposition specification (Column 5), where the composite is replaced by its two LM sub-components, the length coefficient is $+0.130$ ($t = +7.16$). The pattern is consistent across specifications: longer risk-factor sections are associated with higher post-filing log-volatility at the 30-day horizon. The point estimate of 0.130 from Column (5) corresponds to roughly an 8% increase in post-filing volatility for a one-standard-deviation rise in $\ln(\text{Words}^{1A})$ ($0.64 \times 0.130 = 0.083$, applied to $\ln \sigma$). **H1 is supported.**

The specificity channel (Columns 3 and 4). Adding RiskDensity alone (Column 3) yields a coefficient of -7.78 ($t = -5.40$). In the joint specification, where the volume channel is controlled for, the coefficient remains negative and significant at -3.88 ($t = -2.68$, $p = 0.007$). The fact that the density coefficient halves once length is controlled for is exactly the mechanical pattern anticipated in Section 3.2.3: about half of what looked like a “risk-density effect” in Column (3) was in fact the volume effect operating through the section’s word-count denominator. Once $\ln(\text{Words}^{1A})$ is partialled out, the residual density effect is purely specificity in the Hope et al. (2016) sense: a denser, more boilerplate Item 1A is associated with *lower* post-filing volatility. **H2 is supported**, and the joint specification reveals an effect that the dictionary-only ratio used in much prior work would have hidden.

The boilerplate decomposition (Column 5). Replacing the composite RiskDensity by its two LM sub-components reveals that the negative specificity association of Column (4) is carried almost entirely by the LM-Litigious sub-list. LitDensity enters with a coefficient of -6.16 ($t = -3.38$, $p < 0.001$), while UncDensity enters with a small and statistically insignificant coefficient of -1.28 ($t = -0.62$, $p = 0.53$). A formal Wald test of $H_0 : \beta^{\text{unc}} = \beta^{\text{lit}}$ yields $z = +1.91$ ($p = 0.056$), so the two coefficients are distinguishable at the 10% but not the 5% level under the full two-way clustering scheme. The 95% confidence interval on β^{unc} is wide ($[-5.30, +2.74]$): the data are consistent with a modestly negative uncertainty effect, but cannot rule out zero. Adjusted R^2 in Column (5) is essentially identical to that of the composite specification in Column (4) (0.5068 vs. 0.5064), because the litigious sub-list captures the variance that drives the composite. The H2 specificity association is therefore concentrated in the litigious rather than the uncertainty sub-list, consistent with the cautionary, legal-process language that Cazier et al. (2021) and Nelson and Pritchard (2007) link to mature, legally supervised disclosure environments. **The LM-list decomposition is supported in the cross-section**, with the qualifications discussed in Section 4.5 and the firm-FE caveat of Section 5.3.

Joint significance and incremental fit. The joint test of $H_0 : \beta^{\text{len}} = \beta^{\text{unc}} = \beta^{\text{lit}} = 0$ in Column (5) rejects at the 1% level. Adding the three text variables to the baseline raises adjusted R^2 from 0.492 (Column 1) to 0.507 (Column 5), an increment that is small in absolute terms because the financial baseline already captures the bulk of predictable volatility but is

large by the standards of the textual-disclosure literature (Campbell et al., 2014; Kravet & Muslu, 2013). The three text variables are informative both individually and jointly.

4.3 The same channels at the 365-day horizon and across intermediate windows

Table 4.4 re-estimates the same five nested specifications with the dependent variable replaced by $\ln \sigma_{i,t+1}^{[365]}$. The signs and significance of all three text variables are preserved: in the H1+H2 decomposition (Column 5) $\beta^{\text{len}} = +0.061$ ($t = +4.93$), $\beta^{\text{lit}} = -3.39$ ($t = -3.91$), and β^{unc} remains null. Magnitudes shrink to roughly half of the 30-day estimates, consistent with the post-filing reaction documented by Loughran and McDonald (2011).

Table 4.4: Item 1A Length, Risk-Word Density, and the Boilerplate Decomposition (365 days window)

	(1)	(2)	(3)	(4)	(5)
$\ln(\sigma_{i,t})$	+0.6112*** (+14.37)	+0.5731*** (+12.60)	+0.5935*** (+13.41)	+0.5680*** (+12.34)	+0.5657*** (+12.25)
$\ln(\text{Total Assets})$	-0.0339*** (-4.76)	-0.0346*** (-4.94)	-0.0346*** (-4.91)	-0.0349*** (-5.00)	-0.0337*** (-4.76)
Leverage	-0.0116 (-0.81)	-0.0026 (-0.17)	-0.0139 (-0.94)	-0.0051 (-0.34)	-0.0056 (-0.37)
ROA	-0.4402*** (-7.45)	-0.3901*** (-7.20)	-0.4552*** (-7.54)	-0.4054*** (-7.43)	-0.4101*** (-7.42)
Asset Growth	+0.0728** (+2.49)	+0.0586** (+2.29)	+0.0684** (+2.38)	+0.0579** (+2.25)	+0.0568** (+2.20)
$\ln(\text{Words}^{1A})$		+0.0645*** (+5.15)		+0.0558*** (+4.53)	+0.0607*** (+4.93)
RiskDensity			-3.8105*** (-5.42)	-2.1913*** (-3.55)	
UncDensity					-0.8552 (-1.08)
LitDensity					-3.3905*** (-3.91)
Industry FE	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes
Observations	6,172	6,172	6,172	6,172	6,172
Firms	535	535	535	535	535
Adj. R^2	0.729	0.734	0.731	0.735	0.735

Note: The dependent variable is $\ln(\sigma_{i,t+1})$, the natural log of annualised post-filing volatility over a 365 days window starting two trading days after the 10-K filing date. Column (1) reports the financial baseline; column (2) adds the log word count of Item 1A (H1); column (3) adds the composite LM risk-word density (H2); column (4) adds H1 and H2 jointly; column (5) replaces the H2 composite by its LM-list decomposition into the LM-Uncertainty and LM-Litigious densities of Item 1A. All specifications include two-digit SIC industry and fiscal-year fixed effects. t -statistics in parentheses use two-way (firm and fiscal-year) clustered standard errors. ***, **, * denote significance at the 1%, 5%, and 10% levels.

4.4 Multi-horizon table

Table 4.5 reports the H1+H2 decomposition specification of equation 3.9 at six post-filing horizons: $h \in \{5, 10, 30, 90, 180, 365\}$ calendar days. Each column is a separate regression on the full panel; financial controls and fixed effects are unchanged.

Table 4.5: Post-Filing Volatility at Alternative Horizons

	(1) 5d	(2) 10d	(3) 30d	(4) 90d	(5) 180d	(6) 365d
$\ln(\sigma_{i,t})$	+0.1008*** (+3.62)	+0.1249*** (+4.55)	+0.1321* (+1.75)	+0.2122** (+2.24)	+0.3533*** (+4.19)	+0.5657*** (+12.25)
$\ln(\text{Words}^{1A})$	+0.1953*** (+8.78)	+0.1524*** (+6.72)	+0.1298*** (+7.16)	+0.1155*** (+5.93)	+0.0962*** (+5.11)	+0.0607*** (+4.93)
UncDensity	+0.9530 (+0.27)	-2.1555 (-1.04)	-1.2779 (-0.62)	-0.7636 (-0.45)	-0.6946 (-0.55)	-0.8552 (-1.08)
LitDensity	-7.0957*** (-3.76)	-6.6115*** (-2.88)	-6.1591*** (-3.38)	-6.5495*** (-3.94)	-5.1636*** (-4.02)	-3.3905*** (-3.91)
Financial controls	Yes	Yes	Yes	Yes	Yes	Yes
Industry FE	Yes	Yes	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes	Yes	Yes
Observations	5,690	5,648	5,691	5,691	5,691	6,172
Firms	535	535	535	535	535	535
Adj. R^2	0.156	0.354	0.507	0.575	0.643	0.735

Note: Each column re-estimates the H1 + LM-list decomposition specification replacing the dependent variable with $\ln(\sigma_{i,t+1}^{[h]})$ measured over a post-filing window of length h . The lagged dependent variable is the same firm's log-volatility computed over the prior filing year's window of identical length. UncDensity and LitDensity are the LM-Uncertainty and LM-Litigious word densities on Item 1A; together they sum to RiskDensity (H2). Financial controls (size, leverage, ROA, asset growth) are included but suppressed. Industry and year fixed effects are absorbed; t -statistics in parentheses use two-way (firm and fiscal-year) clustered standard errors. ***, **, * denote significance at the 1%, 5%, and 10% levels.

Three patterns emerge. First, $\ln(\text{Words}^{1A})$ enters with a positive and highly significant coefficient at every horizon, ranging from +0.195 at 5 days to +0.061 at 365 days, with t -statistics between +4.93 and +8.78. The volume association is robust across horizons. Second, LitDensity enters with a negative coefficient at every horizon and is significant at the 1% level from 10 days onward; the t -statistics range from -2.88 at 10 days to -4.02 at 180 days. Third, UncDensity is small in absolute value and statistically indistinguishable from zero at every horizon, with $|t| \leq 1.1$ at all windows. The qualitative pattern of Column (5) of Table 4.3 therefore replicates at every horizon the data admit: the negative association between risk-word language and post-filing volatility is a litigious-language association, not an uncertainty-language one.

The lagged-dependent-variable coefficient also rises with horizon, from about 0.030 at 5 days

to 0.566 at 365 days, reflecting the mechanical fact that longer windows for both the dependent and the lagged variable share more of their underlying daily-return sample. This is a feature of the data, not of the specification.

4.5 The boilerplate decomposition

The negative association of the composite RiskDensity with post-filing volatility is driven entirely by the LM-Litigious sub-list; once the litigious component is partialled out, the LM-Uncertainty component has no incremental association at any horizon. A one-standard-deviation rise in LitDensity (≈ 0.005) corresponds to a $0.005 \times 6.16 \approx 2.8\%$ decrease in 30-day annualised volatility, and the interquartile swing is $\approx 3.8\%$. The reading aligns with the Beatty et al. (2019), Cazier et al. (2021), and Nelson and Pritchard (2007) interpretation of litigious-list language as standardised cautionary boilerplate that proxies for an institutionalised, legally supervised compliance function. The within-firm caveat is reserved for Section 5.3.

4.6 Strategic silence and within-firm post-filing volatility

Table 4.6 reports the H3 strategic-silence specification of equation 3.10 at five post-filing horizons under both fixed-effect structures: industry plus year FE in Panel A and firm plus year FE in Panel B. The coefficient of interest is β^{om} on the peer-relative omission gap of equation 3.7. Item 1A length is in the regression as a control alongside RiskDensity, the five financial controls, and lagged log-volatility, so β^{om} measures the incremental association of peer-relative omission with post-filing volatility holding the firm’s own length and risk-vocabulary density fixed.

Cross-section is null. Panel A of Table 4.6 shows that, conditional on the H1+H2 controls and SIC2-and-year fixed effects, peer-relative omission has no detectable cross-sectional association with post-filing volatility at any horizon: $|t| < 1.1$ at every window. The cross-sectional version of equation 3.10 does not identify a between-firm strategic-silence effect on this panel.

Within-firm survival and economic magnitude. Panel B reports the strict version with industry fixed effects replaced by firm fixed effects. The omission coefficient is positive and significant at the 30-day horizon ($\hat{\beta}^{\text{om}} = +1.345$, $t = +2.93$, $p = 0.003$), and remains significant at 90 and 180 days ($t = +2.42$ and $+2.35$); at 5 and 10 days the coefficient is positive but insignificant. A one-within-firm-SD rise in OMISSION^{1A} (0.054) implies a $\approx 7.3\%$ increase in 30-day annualised volatility, of the same order as the cross-sectional H1 length effect. **H3**

Table 4.6: Silence Hypothesis: Item 1A omission relative to SIC2 peers and post-filing realised volatility.

	$h = 5$	$h = 10$	$h = 30$	$h = 90$	$h = 180$
<i>Panel A. Industry $\times$ year fixed effects</i>					
OMISSION _{1A}	+0.424 (+0.71)	-0.054 (-0.11)	+0.098 (+0.28)	+0.179 (+0.63)	+0.226 (+1.04)
N	4,630	4,603	4,636	4,636	4,636
R^2	0.174	0.354	0.505	0.575	0.650
<i>Panel B. Firm $\times$ year fixed effects</i>					
OMISSION _{1A}	+0.808 (+0.81)	+1.233 (+1.50)	+1.345*** (+2.93)	+1.299** (+2.42)	+1.136** (+2.35)
N	4,630	4,603	4,636	4,636	4,636
R^2	0.322	0.484	0.627	0.691	0.750

Notes. Dependent variable is $\log(RV_{hd})$ over the h -day window beginning the trading day after filing.

OMISSION_{1A} is the share of TF-IDF mass on terms used by SIC2-year peers (leave-one-out, ≥ 5 peers) that the focal firm does not mention. All regressions control for $\log(RV_{hd,t-1})$, the five financial controls (size, leverage, ROA, current ratio, OCF/assets), $\log(\text{words}_{1A})$, and RISKDENSITY_{1A}. t -statistics in parentheses use two-way clustered standard errors on (ticker, fiscal year). * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

is supported within-firm at the 30-day horizon, with the qualifications discussed in Section 5.3.

Cross-section vs. within-firm. The H3 omission gap shows the opposite pattern to the H2 LM sub-densities: no cross-sectional power but a within-firm signal. The reading is that permanent firm-level disclosure styles absorb the cross-section, whereas the omission gap, by construction, varies with what *other* firms write each year and is therefore a year-specific peer-relative signal that the firm-FE specification lets through. An alternative reading—omission proxies for a year-specific firm shock that drives both shorter Item 1A and higher subsequent volatility—cannot be ruled out from these regressions alone; the placebo, length-residualised, and lead-lag checks of Section 5.3 are designed to discriminate between the two.

5. Robustness

This chapter checks that the headline results of Chapter 4 survive plausible variations in the specification. Section 5.1 reviews the multi-horizon evidence already documented in Section 4.4 from a robustness angle. Section 5.2 re-estimates the H1+H2 decomposition specification on the pre-COVID and COVID-era sub-samples. Section 5.3 replaces industry fixed effects with firm fixed effects, which absorbs all time-invariant firm-level heterogeneity and isolates the within-firm year-to-year variation in the three text variables. Section 5.3 reports the placebo, length-residualised, large-cell, decile-rank, and prior-year specifications of the H3 strategic-silence result. Section 5.4 closes with a brief note on alternative specificity-channel candidates that were tested and rejected before settling on the LM-list decomposition.

5.1 Alternative post-filing windows

The headline specification measures post-filing volatility over a 30-day window. One month is short enough to capture the bulk of the post-filing reaction documented by Loughran and McDonald (2011) but long enough that the volatility estimate is built on roughly 21 trading days of returns. To check the choice of window we re-estimate the H1+H2 specification of equation 3.9 at five additional horizons (5, 10, 90, 180, and 365 days) and inspect the resulting series of coefficients. Table 4.5 reports all six columns side by side; the discussion in Section 4.4 is repeated here from a robustness angle.

The pattern is consistent in three respects. First, all three text coefficients keep the predicted sign at every horizon—positive for $\ln(\text{Words}^{1A})$, negative for LitDensity, and statistically indistinguishable from zero for UncDensity—so neither H1 nor the LM-list decomposition reverses for any plausible choice of window. Second, the LitDensity coefficient is significant at the 1% level from 10 days onward, with $|t|$ between 2.88 (10 days) and 4.02 (180 days). Third, the UncDensity coefficient never approaches statistical significance, with $|t| \leq 1.1$ at every window. The 30-day choice for the headline specification is therefore not a window in which the results happen to be sharpest—they are sharp at every horizon the data admit.

Both coefficient series also decline approximately monotonically in h , the empirical signature of an event-driven signal that is incorporated into prices over weeks rather than years (Loughran & McDonald, 2011; Tetlock, 2007). No formal half-life is fitted because the across-horizon error terms are correlated through the shared underlying daily returns.

5.2 Pre-COVID and COVID-era sub-samples

The 2020 fiscal year and its aftermath were a structural shock for both equity volatility and corporate disclosure: average annualised volatility roughly doubled in the spring of 2020 and

Item 1A sections grew systematically longer as firms added pandemic and supply-chain language. To check that the headline results are not driven by either the pre-COVID baseline or the COVID-era spike, we re-estimate the H1+H2 decomposition specification (equation 3.9) separately on the two sub-periods. Filings with fiscal year ≤ 2019 form the pre-COVID sample (3,243 firm-years) and filings with fiscal year ≥ 2020 form the COVID-era sample (2,448 firm-years). All other features of the specification are unchanged. Table 5.1 reports the result.

Table 5.1: Sub-Period Robustness (30-day horizon)

	(1) Pre-COVID (2010–2019)	(2) COVID era (2020–2024)
$\ln(\sigma_{i,t})$	+0.1271 (+1.21)	+0.1313* (+1.72)
$\ln(\text{Words}^{1A})$	+0.1025*** (+6.06)	+0.1715*** (+8.92)
UncDensity	−2.8830 (−1.17)	+0.8033 (+0.33)
LitDensity	−5.4863** (−2.55)	−9.8819*** (−4.85)
Financial controls	Yes	Yes
Industry FE	Yes	Yes
Year FE	Yes	Yes
Observations	3,243	2,448
Firms	470	527
Adj. R^2	0.529	0.426

Note: H1 + LM-list decomposition specification re-estimated on two non-overlapping sub-samples. The dependent variable is $\ln(\sigma_{i,t+1}^{[30]})$. UncDensity and LitDensity are the LM-Uncertainty and LM-Litigious word densities on Item 1A. All financial controls and SIC2+year fixed effects are included. Standard errors are two-way clustered by firm and fiscal year. ***, **, * denote significance at the 1%, 5%, and 10% levels.

The two active text channels are present in both sub-periods with the predicted signs and at the same order of magnitude. Length is significant at the 1% level in both ($\beta = +0.103$, $t = +6.06$ pre-COVID; $\beta = +0.172$, $t = +8.92$ COVID era). LitDensity is negative and significant in both ($\beta = -5.49$, $t = -2.55$ pre-COVID; $\beta = -9.88$, $t = -4.85$ COVID era). UncDensity is statistically indistinguishable from zero in both samples, consistent with the headline finding that uncertainty wording adds no explanatory power once length and litigious wording are controlled for. The two active coefficients (length and LitDensity) are larger in absolute value in the COVID-era sample, consistent with the higher overall level of volatility in that period giving more scale to the disclosure signal, but the qualitative pattern is identical. The headline finding is not a COVID artefact.

5.3 Within-firm identification: firm fixed effects

The headline specification absorbs two-digit SIC industry and fiscal-year fixed effects. Coefficients are therefore identified from variation *within an industry-year cell*, which removes the most plausible confounders (industry-level differences in risk vocabulary and macroeconomic time effects) but still allows permanent firm-level traits to drive the result. To rule that out we replace the industry fixed effects with firm fixed effects, so that identification comes only from *within-firm* year-to-year variation in the three text variables. Table 5.2 reports the result at the 30-day horizon.

Table 5.2: Within-Firm Identification (Firm + Year FE, 30-day horizon)

	Estimate
$\ln(\sigma_{i,t})$	−0.1059*** (−6.77)
$\ln(\text{Words}^{1A})$	+0.0835*** (+3.31)
UncDensity	−1.4559 (−0.55)
LitDensity	+0.9561 (+0.35)
Financial controls	Yes
Firm FE	Yes
Year FE	Yes
Observations	5,691
Firms	535
Adj. R^2	0.583

Note: H1 + LM-list decomposition specification with firm and fiscal-year fixed effects absorbed. Identification of $\ln(\text{Words}^{1A})$, UncDensity and LitDensity comes only from within-firm variation across years. Standard errors clustered by firm. ***, **, * denote significance at the 1%, 5%, and 10% levels.

The volume channel survives the firmer specification: $\hat{\beta}^{\text{len}} = +0.084$ ($t = +3.31$, $p < 0.001$). A within-firm one-standard-deviation rise in $\ln(\text{Words}^{1A})$ still predicts a meaningful rise in post-filing volatility, so the H1 result is not a cross-sectional artefact—firms that grow their Item 1A from year to year do experience higher subsequent volatility.

Both LM sub-densities, by contrast, are absorbed by firm fixed effects: $\hat{\beta}^{\text{unc}} = -1.46$ ($t = -0.55$) and $\hat{\beta}^{\text{lit}} = +0.96$ ($t = +0.35$). Risk-vocabulary composition has very little within-firm variation across years because Item 1A is largely recycled from one filing to the next (Cazier et al., 2021); the firm-FE specification therefore has almost no remaining variance to identify a sub-density coefficient from. This is consistent with the production-side reading of Beatty et al. (2019), Cazier et al. (2021), and Nelson and Pritchard (2007), in which cautionary boilerplate is a quasi-permanent firm-level disclosure-style trait. The thesis reports both

findings and reads the headline result as a between-firm regularity, without claiming a within-firm causal channel.

Bottom line. H1 clears all robustness checks including firm-FE; the H2 composite and its LM sub-components clear the cross-sectional and sub-period checks but are absorbed by firm-FE because both RiskDensity and the two LM sub-densities are quasi-permanent firm-level traits.

The headline H3 result of Section 4.6 relies on a within-firm specification at the 30-day horizon. Three features of the construction make further robustness checks essential before reading the result as evidence for strategic silence rather than for some confounding artefact: (i) the omission gap is mechanically negatively correlated with Item 1A length ($r = -0.85$), so any insufficiently strong control for length would let omission proxy for length; (ii) the SIC2 peer benchmark is one of several plausible reference sets, and the result might be specific to that choice; (iii) the contemporaneous design has a partial overlap between the filing date and the start of the post-filing volatility window, leaving some scope for reverse causality if a difficult year produces both a shorter, less peer-aligned 10-K and higher subsequent volatility for unrelated reasons. Table 5.3 reports six robustness specifications designed to address these concerns; a placebo test is reported separately below.

Table 5.3: Robustness of the strategic-silence result (H3): firm $\times$ year fixed effects, $h = 30$ days.

	(1) Baseline	(2) Decile rank	(3) Length-resid	(4) ≥ 20 peers	(5) Triple length	(6) Prior-year
OMISSION _{1A}	+1.345*** (+2.93)	+0.010 (+1.31)	+0.352 (+1.49)	+2.519*** (+3.09)	+1.502*** (+3.18)	+0.485** (+1.97)
N	4,636	3,706	4,636	3,248	4,636	4,546

Notes. Each column is a separate firm $\times$ year fixed-effect regression at $h = 30$ days. Column (1) reports the baseline H3 specification of equation 3.10. Column (2) replaces OMISSION_{1A} with its decile rank within (SIC2, fiscal year) cells. Column (3) replaces it with the residual after regressing OMISSION_{1A} on $\ln(\text{WORDS}_{1A})$ within (SIC2, fiscal year) cells. Column (4) restricts the sample to industry-year cells with at least 20 panel firms. Column (5) adds $\ln(\text{WORDS}_{MDA})$ and $\ln(\text{WORDS}_{1A})^2$ as additional length controls. Column (6) replaces the contemporaneous omission with its prior-year value. All regressions include lagged log-volatility, the five financial controls, $\ln(\text{WORDS}_{1A})$, and RISKDENSITY_{1A}. t -statistics in parentheses use two-way clustered standard errors on (ticker, fiscal year). * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Placebo (peer-cell shuffle). The first and most important robustness check holds the SIC2 cell sizes and the firm-year distribution constant but randomly permutes which firm belongs to which SIC2 cell within fiscal year, then recomputes the omission gap of equation 3.7 and re-estimates equation 3.10 with firm and year fixed effects. If the omission–volatility association reflects something specifically about being silent on *peer-relevant* risks, the placebo coefficient should die. If the result is a length artefact, the placebo should mirror the real coefficient. Across five independent permutations the placebo coefficient is small and statisti-

cally insignificant in every draw, with $|t|$ values 0.88, 0.90, 0.54, 0.57, 0.63. The real coefficient $t = +2.93$ exceeds the maximum placebo $|t|$ by a factor of three. The strategic-silence reading of the H3 result therefore survives a clean placebo test on the peer-structure mechanism.

Length is not the driver (Columns 3 and 5). Column (3) replaces the omission gap by its residual after regressing on $\ln(\text{WORDS}_{1A})$ within each (SIC2, fiscal year) cell; the purified variation enters with the right sign but loses significance ($t = +1.49$). Column (5) keeps the original gap but adds two further length controls (MD&A word count and squared Item 1A log word count); the omission coefficient strengthens to $t = +3.18$ ($p = 0.001$). Linear or quadratic length controls do not remove the signal, but the strict orthogonalisation of Column (3) absorbs enough of the effect to leave a borderline rather than a strong association.

Large-cell restriction (Column 4) and prior-year specification (Column 6). Restricting to cells with at least 20 peers reduces the panel to 3,248 firm-years but *strengthens* the coefficient to $+2.52$ ($t = +3.09$): the signal is sharper where the peer aggregate is more reliably estimated. Column (6) replaces the contemporaneous omission with its prior-year value to address the filing-window overlap concern; the lagged coefficient remains positive and significant at the 5% level ($+0.485$, $t = +1.97$).

Decile rank (Column 2). Replacing the level of omission by its decile rank within each (SIC2, fiscal year) cell preserves the predicted positive sign but reduces significance to $t = +1.31$. The H3 effect is concentrated in the tails of the omission distribution, consistent with a strategic-silence reading in which conspicuous omissions drive the response. The decile-rank weakness is the H3 result's most important qualification.

Bottom line on H3. The strategic-silence result survives the placebo, two stronger length controls, the large-cell restriction, and the one-year-lag specification, with the right sign in every column. Two qualifications apply: the decile-rank specification is not significant at the 5% level, and the fully length-orthogonalised reformulation (Column 3) is borderline rather than significant. The thesis reads the H3 evidence as supportive of a within-firm strategic-silence channel with these stated qualifications, and does not extend the claim to a causal interpretation.

5.4 What did not work

Seven alternative specificity-channel candidates were estimated on the same panel and rejected before the LM-list decomposition was adopted: (i) year-on-year TF-IDF similarity of Item 1A in the spirit of Brown and Tucker (2011); (ii) MD&A LM-Negativity; (iii) the first difference $\Delta \ln(\text{Words}^{1A})$ following Lyle et al. (2023); and (iv)–(vii) four functional forms of

FinBERT-tone (`yyianghkust/finbert-tone`) aggregated from the per-sentence probabilities ($P(\text{neg})$, the signed $P(\text{pos}) - P(\text{neg})$, $P(\text{neutral})$, and the intensity $1 - P(\text{neutral})$). Each candidate was added as a third text variable to the H1+H2 joint specification of equation 3.8; none added statistically significant explanatory power for $\ln \sigma_{i,t+1}^{[30]}$ ($|t| < 1.5$ in every case). The TF-IDF and FinBERT signals capture firm-specific narrative changes whose informational content is, on this panel, dominated by the much stronger length signal of H1. The LM-list decomposition succeeds where these alternatives fail because it identifies a structural composition of the H2 specificity effect rather than an additional firm-specific signal that competes with H1.

Conclusion

References

- Agresti, A., & Coull, B. A. (1998). Approximate is better than “exact” for interval estimation of binomial proportions. *The American Statistician*, 52(2), 119–126.
- Beatty, A., Cheng, L., & Zhang, H. (2019). Are risk factor disclosures still relevant? evidence from market reactions to risk factor disclosures before and after the financial crisis. *Contemporary Accounting Research*, 36(2), 805–838.
- Blankespoor, E., deHaan, E., & Marinovic, I. (2020). Disclosure processing costs, investors’ information choice, and equity market outcomes: A review. *Journal of Accounting and Economics*, 70(2–3), 101344.
- Brown, S. V., & Tucker, J. W. (2011). Large-sample evidence on firms’ year-over-year MD&A modifications. *Journal of Accounting Research*, 49(2), 309–346.
- Cameron, A. C., Gelbach, J. B., & Miller, D. L. (2011). Robust inference with multiway clustering. *Journal of Business & Economic Statistics*, 29(2), 238–249.
- Campbell, J. L., Chen, H., Dhaliwal, D. S., Lu, H.-m., & Steele, L. B. (2014). The information content of mandatory risk factor disclosures in corporate filings. *Review of Accounting Studies*, 19(1), 396–455.
- Cazier, R. A., McMullin, J. L., & Treu, J. S. (2021). Are lengthy and boilerplate risk factor disclosures inadequate? an examination of judicial and regulatory assessments of risk factor language. *The Accounting Review*, 96(4), 131–155.
- Chychyla, R., & Kogan, A. (2015). Using XBRL to conduct a large-scale study of discrepancies between the accounting numbers in Compustat and SEC 10-K filings. *Journal of Information Systems*, 29(1), 37–72. <https://doi.org/10.2308/isys-50922>
- Dye, R. A. (1985). Disclosure of nonproprietary information. *Journal of Accounting Research*, 23(1), 123–145.
- Dyer, T., Lang, M., & Stice-Lawrence, L. (2017). The evolution of 10-K textual disclosure: Evidence from Latent Dirichlet Allocation. *Journal of Accounting and Economics*, 64(2–3), 221–245.
- Francis, J., LaFond, R., Olsson, P. M., & Schipper, K. (2005). The market pricing of accruals quality. *Journal of Accounting and Economics*, 39(2), 295–327.
- Gilardi, F., Alizadeh, M., & Kubli, M. (2023). ChatGPT outperforms crowd-workers for text-annotation tasks. *Proceedings of the National Academy of Sciences*, 120(30), e2305016120.
- Hope, O.-K., Hu, D., & Lu, H. (2016). The benefits of specific risk-factor disclosures. *Review of Accounting Studies*, 21(4), 1005–1045.
- Kravet, T., & Muslu, V. (2013). Textual risk disclosures and investors’ risk perceptions. *Review of Accounting Studies*, 18(4), 1088–1122.
- Krippendorff, K. (2019). *Content analysis: An introduction to its methodology* (4th). SAGE Publications.
- Li, F. (2008). Annual report readability, current earnings, and earnings persistence. *Journal of Accounting and Economics*, 45(2–3), 221–247.

- Loughran, T., & McDonald, B. (2011). When is a liability not a liability? Textual analysis, dictionaries, and 10-Ks. *The Journal of Finance*, 66(1), 35–65.
- Lyle, M. R., Riedl, E. J., & Siano, F. (2023). Changes in risk factor disclosures and the variance risk premium. *The Accounting Review*, 98(5), 327–360.
- Nelson, K. K., & Pritchard, A. C. (2007). Litigation risk and voluntary disclosure: The use of meaningful cautionary language. *Working Paper, University of Michigan / Rice University*.
- Petersen, M. A. (2009). Estimating standard errors in finance panel data sets: Comparing approaches. *The Review of Financial Studies*, 22(1), 435–480.
- Sloan, R. G. (1996). Do stock prices fully reflect information in accruals and cash flows about future earnings? *The Accounting Review*, 71(3), 289–315.
- Tetlock, P. C. (2007). Giving content to investor sentiment: The role of media in the stock market. *The Journal of Finance*, 62(3), 1139–1168.
- Verrecchia, R. E. (1983). Discretionary disclosure. *Journal of Accounting and Economics*, 5, 179–194.

Appendices

A. Variable Definitions and Data Pipeline

This appendix collects supplementary tables referenced in the main text. Section A.1 presents the diagnostics behind the choice of annual filing frequency (Section 2.2). Section A.2 reports the per-firm XBRL tag assignments from the tag-harmonisation procedure (Section 2.3).

A.1 Filing Frequency Diagnostics

The following tables supplement the discussion in Section 2.2.

Table A.1: XBRL Financial Metric Coverage: Annual (10-K) vs. Quarterly (10-Q)

Metric	Annual (%)	Quarterly (%)
assets_current	87.0	91.1
eps_diluted	89.4	92.7
liabilities_current	87.0	91.1
net_income	93.6	97.3
operating_cash_flow	89.2	35.8
operating_income	76.9	80.5
stockholders_equity	98.7	94.9
total_assets	89.9	93.3
total_liabilities	63.7	66.5
Median periods per firm	20	52
Firms with data	30	30

Note: Coverage shows the percentage of firm-periods with non-missing values for each metric. Annual data is extracted from 10-K filings; quarterly data from 10-Q filings (Q1–Q3 only). Sample: 30 non-financial operating companies from the top 200 by market capitalisation.

Table A.2: Annual vs. Quarterly Approach: Summary Comparison

Criterion	Annual (10-K)	Quarterly (10-Q)
Filing type	10-K	10-Q (Q1–Q3 only)
MD&A availability	93% (Item 7)	100% (Part I, Item 2)
Risk Factors availability	100% (Item 1A)	~52% substantive
Risk Factors content	Median ~30,000 words	Median ~9,000 words (substantive); 48% boilerplate
Q4 financial data	Included in 10-K	Unreliable / absent
Explicit Q4 XBRL facts	N/A	3–43% of firm-years
Implied Q4 mismatch rate	N/A	~11% of comparisons
Cash-flow quarterly coverage	N/A	7% of firm-years
Item 1A length (H1)	Feasible (annual word count)	Limited (52% substantive)
Item 1A risk-word density (H2)	Feasible (annual LM ratio)	Limited (52% substantive)
Literature alignment	Standard	Uncommon
Expected observations	~6,400 firm-years	~19,000 firm-quarters

Note: Summary based on diagnostic tests of 30 non-financial companies (XBRL, Q4 tests) and 15 companies (textual tests). Risk Factors in 10-Q filings are required by SEC Regulation S-K Item 1A only when there are “material changes” from the most recent 10-K; roughly half of companies include substantive risk disclosures, while the rest provide only a boilerplate reference. See Tables A.1, 2.1, and 2.2 for detailed results.

Table A.3: Per-Company 10-Q Textual Content (Most Recent 6 Filings)

Ticker	Filings	MD&A Found	MD&A Med. Words	Risk Subst.	Risk Boilerplate
AAPL	6	6	2,781	3	3
AMZN	6	6	6,291	6	0
AVGO	6	6	4,778	6	0
COST	6	6	4,979	0	6
GOOGL	6	6	8,096	3	3
JNJ	6	6	7,587	0	6
LLY	6	6	4,994	0	6
MA	6	6	7,300	6	0
META	6	6	7,756	6	0
MSFT	6	6	7,981	0	6
NVDA	6	6	5,156	6	0
TSLA	6	6	5,582	2	4
V	6	6	5,043	0	6
WMT	6	6	7,142	3	3
XOM	6	6	6,197	6	0

Note: “MD&A Found” and “Risk Subst.” count filings with >200 substantive words. “Risk Boilerplate” counts filings where the risk-factors section contains only a reference to the annual 10-K.

A.2 XBRL Tag Provenance

Table A.4 reports, for each metric with multiple candidate XBRL tags, the number and share of firms locked to each tag by the per-firm tag-locking procedure described in Section 2.3.

Table A.4: Per-firm XBRL tag distribution

Metric	Locked XBRL tag	Firms	%
Net income	NetIncomeLoss	502	91.6
	ProfitLoss	45	8.2
	IncomeLossFromContinuingOperations	1	0.2
Operating cash flow	NetCashProvidedByUsedInOperatingActivities	530	96.7
	...OperatingActivitiesContinuingOperations	18	3.3
Stockholders' equity	StockholdersEquityIncl...NoncontrollingInterest	318	58.6
	StockholdersEquity	225	41.4

Notes. For each (firm, metric) pair, the tag with the highest within-firm year-count is selected; values from other tags are set to missing. Firm counts are based on up to 548 firms in the tag-locking stage; the stockholders' equity row sums to 543 because five firms report no equity tag at all. See Section 2.3 for the rationale.

B. Documentation of AI Tool Usage

This appendix documents the use of artificial intelligence tools in preparing this thesis, as required by the university.

Extraction Validation (Section 2.4.4)

Tool. Google Gemini 2.5 Flash (accessed via the OpenAI-compatible Gemini API endpoint), with temperature set to 0 for deterministic output.

Purpose. An independent, automated quality audit of the SEC 10-K section extraction pipeline. The model evaluated 200 extracted sections (100 Item 1A, 100 Item 7) against a structured five-criterion rubric (correct section, start boundary, end boundary, content quality, completeness) on a three-level scale (Pass / Minor / Fail).

Contribution and verification. The evaluation rubric, sampling strategy, prompt design, and statistical analysis (Wilson confidence intervals) were all designed by the author. The model’s only role was applying the rubric to each extracted section. All Fail verdicts were manually spot-checked against the original SEC filings and confirmed to be correct. Results are reported in Table B.1.

Table B.1: LLM-assisted extraction validation results ($n = 200$ sections, 100 per section type).

Criterion	Strict	%	Lenient	%	95 % CI (lenient)
Correct section	192/200	96.0	192/200	96.0	[92.3 – 98.0 %]
Start boundary	191/200	95.5	193/200	96.5	[93.0 – 98.3 %]
End boundary	192/200	96.0	193/200	96.5	[93.0 – 98.3 %]
Content quality	191/200	95.5	195/200	97.5	[94.3 – 98.9 %]
Completeness	189/200	94.5	191/200	95.5	[91.7 – 97.6 %]
Overall (all five)	179/200	89.5	187/200	93.5	[89.2 – 96.2 %]

Code Development

Tool. AI-powered code completion (integrated in VS Code).

Purpose. Code completion during the development of the data-pipeline scripts that build the panel and extract the 10-K text sections.

Contribution and verification. All generated code was reviewed, tested, and modified by the author. The overall architecture, variable definitions, and analytical choices are the author's own.

Writing Assistance

Tool. AI-powered writing assistant (integrated in VS Code).

Purpose. Grammar and style suggestions for English-language prose in the thesis text.

Contribution and verification. All substantive content, argumentation, and interpretation are the author's own. The tool was used only for language polishing; every suggestion was reviewed before acceptance.